

N-Gram Language Modeling with Unknown Words

Assignment field		Details
Group Number	90	
Problem Statement	9 - N-gram Language Modeling with Unknown Words	
Dataset	NLTK Gutenberg Corpus	
Member 1	HARI KRISHNA MOORTHY T - 2025ag05393	
Member 2	SHREEJA MAHAPATRA - 2025ag05424	
Member 3	NAVEEN RANJOLKAR - 2025ag0539	
Member 4	SOMVANSHI PRASAD DIGAMBAR - 2025ag05285	

1. Introduction

A **language model** assigns probabilities to sequences of words and can use them to predict likely words. An **N-gram** is a sequence of N consecutive tokens. A **trigram** model uses the previous two words to estimate the next word:

$$P(w_i | w_{i-2}, w_{i-1})$$

Words absent from the training vocabulary are a problem because the model has no counts for them. We use the special token `<UNK>` (unknown) to group rare training words and represent unseen test words. This gives the model a learned way to handle unfamiliar vocabulary.

2. Objective

This assignment will:

- build Corpus preprocessing;
- build a trigram language model;

- analyze rare words and replace training words occurring fewer than 3 times with `<UNK>` ;
- perform next-word and missing-word prediction;
- compare models with and without `<UNK>` ;
- compare trigram models with and without `<UNK>` handling; and
- evaluate accuracy, OOV rate, coverage, and perplexity.

3. Import libraries and reproducible setup

The setup checks for NLTK resources and downloads them only when absent. The fixed seed makes any sampling deterministic.

In [184...

```
# Standard-library tools for paths, counting, mathematics, randomness, and text matching.
from pathlib import Path
from collections import Counter
import math
import random
import re
import sys

# Third-party libraries used for plots, corpus access, arrays, and tables.
import matplotlib.pyplot as plt
import nltk
import numpy as np
import pandas as pd
from IPython.display import display

# One fixed seed makes the random train/test split reproducible.
SEED = 42
random.seed(SEED)
np.random.seed(SEED)
pd.set_option("display.max_colwidth", 90)

# Store NLTK resources inside the project so the corpus is available locally.
# The folder already contains the official NLTK packages, but the download loop
# also makes the notebook work if a package is accidentally removed.
NLTK_DATA_DIR = Path.cwd() / "data" / "nltk_data"
NLTK_DATA_DIR.mkdir(parents=True, exist_ok=True)

# Search the project-local folder before checking NLTK's other data locations.
if str(NLTK_DATA_DIR) not in nltk.data.path:
    nltk.data.path.insert(0, str(NLTK_DATA_DIR))

# Map each package name to the path that NLTK expects after installation.
required_resources = {
    "gutenberg": "corpora/gutenberg.zip",
    "punkt": "tokenizers/punkt.zip",
    "punkt_tab": "tokenizers/punkt_tab.zip",
}

# Download only resources that are missing; existing local files are reused.
for package, resource_path in required_resources.items():
```

```

try:
    nltk.data.find(resource_path)
except LookupError:
    nltk.download(
        package,
        download_dir=str(NLTK_DATA_DIR),
        quiet=True,
        raise_on_error=True,
    )

from nltk.corpus import gutenberg

print(f"Python: {sys.version.split()[0]}")
print(f"NLTK: {nltk.__version__}")
print(f"Random seed: {SEED}")
print(f"Local NLTK data folder: {NLTK_DATA_DIR.resolve()}")

```

Python: 3.14.6

NLTK: 3.10.3

Random seed: 42

Local NLTK data folder: /Users/harikrishnamoorthy/github/ml/bits-assignment/gutenberg_nlp/Group90_GutenbergDataset/data/nltk_data

4. Load and inspect the NLTK Gutenberg Corpus

We first list every available Gutenberg file and then load all books. Combining the complete corpus provides a larger and more varied collection of sentences for training, rare-word analysis, and `<UNK>` evaluation.

In [185...

```

# Ask NLTK for every file identifier available in its Gutenberg Corpus.
available_books = gutenberg.fileids()

# Convert the list to a small table so it is easy to read in the notebook.
books_df = pd.DataFrame({"Available Gutenberg file": available_books})
display(books_df)
print(f"Number of books available: {len(available_books)}")
print("Selection: all available Gutenberg books")

```

	Available Gutenberg file
0	austen-emma.txt
1	austen-persuasion.txt
2	austen-sense.txt
3	bible-kjv.txt
4	blake-poems.txt
5	bryant-stories.txt
6	burgess-busterbrown.txt
7	carroll-alice.txt
8	chesterton-ball.txt
9	chesterton-brown.txt
10	chesterton-thursday.txt
11	edgeworth-parents.txt
12	melville-moby_dick.txt
13	milton-paradise.txt
14	shakespeare-caesar.txt
15	shakespeare-hamlet.txt
16	shakespeare-macbeth.txt
17	whitman-leaves.txt

Number of books available: 18

Selection: all available Gutenberg books

In [186...

```
# These containers collect the text, sentences, tokens, and statistics from all books.
raw_text_parts, raw_sentences, raw_tokens = [], [], []
book_statistics = []

# Read each book through NLTK. No external text file is used.
for book_id in available_books:
    book_raw_text = gutenberg.raw(book_id)
    book_sentences = list(gutenberg.sents(book_id))
    book_tokens = list(gutenberg.words(book_id))

    # append() keeps one raw-text item per book; extend() creates flat sentence/token lists.
    raw_text_parts.append(book_raw_text)
    raw_sentences.extend(book_sentences)
    raw_tokens.extend(book_tokens)
    book_statistics.append({
        "Book": book_id,
        "Raw characters": len(book_raw_text),
        "Raw sentences": len(book_sentences),
        "Raw tokens": len(book_tokens),
    })

# One combined string is convenient for total character statistics.
raw_text = "\n".join(raw_text_parts)

# Summarize the complete combined corpus.
corpus_stats_df = pd.DataFrame({
    "Statistic": ["Books loaded", "Raw characters", "Raw sentences", "Tokens before preprocessing"],
    "Value": [len(available_books), sum(len(text) for text in raw_text_parts),
              len(raw_sentences), len(raw_tokens)],
})
display(corpus_stats_df)
```

```
per_book_stats_df = pd.DataFrame(book_statistics)
display(per_book_stats_df)
```

		Statistic	Value
	0	Books loaded	18
	1	Raw characters	11793318
	2	Raw sentences	98552
	3	Tokens before preprocessing	2621613

	Book	Raw characters	Raw sentences	Raw tokens
0	austen-emma.txt	887071	7752	192427
1	austen-persuasion.txt	466292	3747	98171
2	austen-sense.txt	673022	4999	141576
3	bible-kjv.txt	4332554	30103	1010654
4	blake-poems.txt	38153	438	8354
5	bryant-stories.txt	249439	2863	55563
6	burgess-busterbrown.txt	84663	1054	18963
7	carroll-alice.txt	144395	1703	34110
8	chesterton-ball.txt	457450	4779	96996
9	chesterton-brown.txt	406629	3806	86063
10	chesterton-thursday.txt	320525	3742	69213
11	edgeworth-parents.txt	935158	10230	210663
12	melville-moby_dick.txt	1242990	10059	260819
13	milton-paradise.txt	468220	1851	96825
14	shakespeare-caesar.txt	112310	2163	25833
15	shakespeare-hamlet.txt	162881	3106	37360
16	shakespeare-macbeth.txt	100351	1907	23140
17	whitman-leaves.txt	711215	4250	154883

5. Text preprocessing

Preprocessing decisions are explicit:

1. NLTK's Gutenberg sentence boundaries are retained.
2. Tokens are lowercased so `The` and `the` share counts.
3. Alphabetic words and internal apostrophes are kept; standalone punctuation and numbers are removed.
4. Two `<START>` tokens and one `<END>` token are added during trigram counting. Two starts are required because a trigram needs two preceding tokens.
5. `<UNK>` is **not** added yet; it is learned from training frequencies later.

```
In [187... # Accept alphabetic words and words with an internal apostrophe; reject punctuation/numbers.
WORD_PATTERN = re.compile(r"^[A-Za-z]+(?:'[A-Za-z]+)?$")
```

```
def preprocess_sentence(tokens):
    # Lowercase and retain word-like tokens while preserving the sentence as a list.
    return [token.lower() for token in tokens if WORD_PATTERN.fullmatch(token)]

# Process one sentence at a time so sentence boundaries are not lost.
processed_sentences = [preprocess_sentence(sentence) for sentence in raw_sentences]
processed_sentences = [sentence for sentence in processed_sentences if sentence]

# Show before/after examples to make every preprocessing decision visible.
example_rows = []
for original, processed in zip(raw_sentences[:5], processed_sentences[:5]):
    example_rows.append({
        "Original sentence": " ".join(original[:30]),
        "After preprocessing": " ".join(processed[:30]),
    })
display(pd.DataFrame(example_rows))
```

	Original sentence	After preprocessing
0	[Emma by Jane Austen 1816]	emma by jane austen
1	VOLUME I	volume i
2	CHAPTER I	chapter i
3	Emma Woodhouse , handsome , clever , and rich , with a comfortable home and happy disp...	emma woodhouse handsome clever and rich with a comfortable home and happy disposition ...
4	She was the youngest of the two daughters of a most affectionate , indulgent father ; ...	she was the youngest of the two daughters of a most affectionate indulgent father and ...

6. Train/test data preparation

The corpus is randomly split at sentence boundaries: 80% of complete sentences are used for training and 20% are used for testing. The fixed random seed makes the split reproducible. The split is completed before vocabulary and `<UNK>` statistics are learned, which prevents test information from leaking into training.

```
In [188... # Shuffle sentence indices rather than individual words so sentence boundaries remain intact.
sentence_indices = list(range(len(processed_sentences)))
split_random = random.Random(SEED)
split_random.shuffle(sentence_indices)

# Use 80% of the shuffled sentence indices for training and 20% for testing.
split_index = int(0.80 * len(sentence_indices))
train_indices = sentence_indices[:split_index]
test_indices = sentence_indices[split_index:]

# Select whole sentences; words from one sentence never cross between the two splits.
train_sentences = [processed_sentences[index] for index in train_indices]
test_sentences = [processed_sentences[index] for index in test_indices]

split_df = pd.DataFrame({
    "Split": ["Training", "Testing", "Total"],
    "Sentences": [len(train_sentences), len(test_sentences), len(processed_sentences)],
    "Percentage": [100 * len(train_sentences) / len(processed_sentences),
                  100 * len(test_sentences) / len(processed_sentences), 100.0],
})
```

```
})
display(split_df.style.format({"Percentage": "{:.1f}%"})
```

	Split	Sentences	Percentage
0	Training	78656	80.0%
1	Testing	19665	20.0%
2	Total	98321	100.0%

7. Exploratory Data Analysis

EDA is reported on the training split because training frequencies determine the vocabulary and rare-word rule. This also avoids looking at test data when designing `<UNK>` .

In [189...

```
# Flatten each split only for frequency calculations; the sentence lists remain unchanged.
train_tokens = [word for sentence in train_sentences for word in sentence]
test_tokens = [word for sentence in test_sentences for word in sentence]
train_frequency = Counter(train_tokens)

# Count vocabulary types with frequencies 1, 2, and below the required threshold of 3.
words_once = sum(count == 1 for count in train_frequency.values())
words_twice = sum(count == 2 for count in train_frequency.values())
rare_types = sum(count < 3 for count in train_frequency.values())
rare_token_count = sum(count for count in train_frequency.values() if count < 3)


frequency_stats_df = pd.DataFrame({
    "Statistic": ["Total training tokens", "Unique training words", "Words occurring once",
                 "Words occurring twice", "Word types occurring < 3 times",
                 "Rare-token occurrences", "Rare types as % of vocabulary"],
    "Value": [len(train_tokens), len(train_frequency), words_once, words_twice,
              rare_types, rare_token_count, 100 * rare_types / len(train_frequency)],
})
display(frequency_stats_df.style.format({"Value": lambda x: f"{x:,.2f}" if isinstance(x, float) else f"{x:,}"})
display(pd.DataFrame(train_frequency.most_common(20), columns=["Word", "Frequency"])))
```

	Statistic	Value
0	Total training tokens	1,708,115.00
1	Unique training words	37,932.00
2	Words occurring once	13,943.00
3	Words occurring twice	5,299.00
4	Word types occurring < 3 times	19,242.00
5	Rare-token occurrences	24,541.00
6	Rare types as % of vocabulary	50.73

	Word	Frequency
0	the	106914
1	and	76394
2	of	57045
3	to	38361
4	a	27247
5	in	26856
6	i	24249
7	that	23090
8	he	20715
9	it	17934
10	his	17096
11	for	15583
12	was	14989
13	with	14044
14	not	13888
15	is	13172
16	you	13070
17	be	12934
18	as	11714
19	but	11134

In [190...


```
# Plot the twenty words with the largest training counts.
top_words = pd.DataFrame(train_frequency.most_common(20), columns=["Word", "Frequency"])
fig, ax = plt.subplots(figsize=(11, 5))
ax.bar(top_words["Word"], top_words["Frequency"], color="#4C78A8")
ax.set_title("Top 20 Most Frequent Training Words")
ax.set_xlabel("Word")
ax.set_ylabel("Frequency")
ax.tick_params(axis="x", rotation=55)
plt.tight_layout()
plt.show()
print("Interpretation: Function words dominate the corpus, so they are likely predictions in many common contexts.")
```



Interpretation: Function words dominate the corpus, so they are likely predictions in many common contexts.


```
In [191... # Count how many word types occur once, twice, three times, and so on.
frequency_of_frequencies = Counter(train_frequency.values())
freq_df = pd.DataFrame(sorted(frequency_of_frequencies.items()), columns=["Word frequency", "Number of word types"])

fig, ax = plt.subplots(figsize=(9, 5))
ax.scatter(freq_df["Word frequency"], freq_df["Number of word types"], alpha=0.65, color="#F58518")
# Log scales make both rare and very frequent words visible on one graph.
ax.set_xscale("log")
ax.set_yscale("log")
ax.set_title("Word-Frequency Distribution (Log-Log Scale)")
ax.set_xlabel("Frequency of a word (log scale)")
ax.set_ylabel("Number of word types (log scale)")
ax.grid(alpha=0.2)
plt.tight_layout()
plt.show()
print("Interpretation: Most word types occur only a few times, while a small number occur very often.")
```



Interpretation: Most word types occur only a few times, while a small number occur very often.

```
In [192... # Compare the two groups that will be replaced because their frequency is below 3.
rare_distribution = pd.DataFrame({
    "Training frequency": ["Once", "Twice"],
    "Number of word types": [words_once, words_twice],
})
fig, ax = plt.subplots(figsize=(7, 4))
bars = ax.bar(rare_distribution["Training frequency"], rare_distribution["Number of word types"], color=["#E45756", "#72B7B2"])
ax.bar_label(bars, fmt="%d")
ax.set_title("Rare-Word Frequency Distribution")
ax.set_xlabel("Occurrence count")
ax.set_ylabel("Number of word types")
plt.tight_layout()
plt.show()
print("Interpretation: The large rare-word tail motivates grouping frequency-1 and frequency-2 words as <UNK>.")
```



Interpretation: The large rare-word tail motivates grouping frequency-1 and frequency-2 words as <UNK>.

8. Trigram model without <UNK>

For context $h=(w_{i-2}, w_{i-1})$, the maximum-likelihood estimate is:

$$P(w_i | h) = \frac{\text{Count}(h, w_i)}{\text{Count}(h)}$$

The class below stores trigram, context, and unigram counts. Prediction is deterministic: the highest probability wins, then training frequency and alphabetical order break ties.

```
In [193... # Special markers represent sentence boundaries and unknown words.
START, END, UNK = "<START>", "<END>", "<UNK>"

class TrigramLanguageModel:
    """A simple count-based trigram model with optional <UNK> handling."""
    def __init__(self, sentences, use_unk=False, rare_threshold=3):
        self.use_unk = use_unk
```

```

self.rare_threshold = rare_threshold
# Training counts alone determine the vocabulary and rare words.
original_counts = Counter(word for sentence in sentences for word in sentence)
self.original_vocabulary = set(original_counts)
self.rare_words = {word for word, count in original_counts.items() if count < rare_threshold} if use_unk else set()
self.sentences = [self.transform_sentence(sentence, training=True) for sentence in sentences]
self.word_counts = Counter(word for sentence in self.sentences for word in sentence)
self.vocabulary = set(self.word_counts) | {END}
if use_unk:
    self.vocabulary.add(UNK)
self.prediction_vocabulary = sorted(self.vocabulary - {START})
# Store Count(w1, w2, w3) and Count(w1, w2).
self.trigram_counts = Counter()
self.context_counts = Counter()
for sentence in self.sentences:
    # Two starts supply a complete context for the first predicted word.
    padded = [START, START] + sentence + [END]
    for i in range(2, len(padded)):
        context = (padded[i - 2], padded[i - 1])
        target = padded[i]
        self.trigram_counts[(context[0], context[1], target)] += 1
        self.context_counts[context] += 1

def map_word(self, word):
    # At test time, map words outside the learned UNK vocabulary to <UNK>.
    if not self.use_unk:
        return word
    return word if word in self.vocabulary else UNK

def transform_sentence(self, sentence, training=False):
    # During training replace rare words; during testing replace unseen words.
    if self.use_unk:
        if training:
            return [UNK if word in self.rare_words else word for word in sentence]
        return [self.map_word(word) for word in sentence]
    return list(sentence)

def probability(self, word, context):
    # Maximum-likelihood probability: trigram count divided by context count.
    context = tuple(self.map_word(w) for w in context[-2:])
    word = self.map_word(word)
    if word not in self.vocabulary:
        return 0.0
    count = self.trigram_counts[(context[0], context[1], word)]
    context_count = self.context_counts[context]
    return count / context_count if context_count else 0.0

def predict_next(self, context_words):
    # Use only the final two words because this is a trigram model.
    mapped = [self.map_word(w) for w in context_words]
    context = ([START, START] + mapped)[-2:]
    observed = [w for w in self.prediction_vocabulary
                 if self.trigram_counts[(context[0], context[1], w)] > 0]
    if not observed:
        return None
    return min(observed, key=lambda w: (-self.probability(w, context), -self.word_counts[w], w))

def sentence_perplexity(self, sentences):
    # Return infinity immediately if any held-out trigram has probability zero.

```

```

log_probability_sum = 0.0
predicted_tokens = 0
for sentence in sentences:
    mapped = self.transform_sentence(sentence, training=False)
    padded = [START, START] + mapped + [END]
    for i in range(2, len(padded)):
        probability = self.probability(padded[i], padded[i-2:i])
        predicted_tokens += 1
        if probability <= 0:
            return math.inf
        log_probability_sum += math.log(probability)
    return math.exp(-log_probability_sum / predicted_tokens)

base_model = TrigramLanguageModel(train_sentences, use_unk=False)
print(f"Base vocabulary size (including <END>): {len(base_model.vocabulary):,}")
print(f"Distinct observed trigrams: {len(base_model.trigram_counts):,}")

```

Base vocabulary size (including <END>): 37,933
 Distinct observed trigrams: 1,114,721

9. Trigram model with <UNK>

Only **training frequencies** decide which words are rare. Every training word with frequency `< 3` is replaced by `<UNK>`, and test words outside the resulting training vocabulary are mapped only when passed to this model. Original test sentences remain unchanged.

- A **rare word** has a low training count (here 1 or 2).
- An **unknown word** is unavailable to the model at prediction time.
- An **out-of-vocabulary (OOV) word** is absent from the model's training vocabulary.

In [194...


```

# Rebuild the model after replacing training words with frequency below 3 by <UNK>.
unk_model = TrigramLanguageModel(train_sentences, use_unk=True, rare_threshold=3)
# Compare vocabulary size and the amount of training text represented by <UNK>.
vocabulary_comparison_df = pd.DataFrame({
    "Statistic": ["Vocabulary size", "Rare word types replaced", "Training token occurrences mapped to <UNK>", "Percent of training tokens mapped"],
    "Before <UNK>": [len(base_model.vocabulary), 0, 0, 0.0],
    "After <UNK>": [len(unk_model.vocabulary), len(unk_model.rare_words), unk_model.word_counts[UNK],
                    100 * unk_model.word_counts[UNK] / len(train_tokens)],
})
display(vocabulary_comparison_df.style.format({"Before <UNK>": "{:,.2f}", "After <UNK>": "{:,.2f}"}))

fig, ax = plt.subplots(figsize=(7, 4))
values = [len(base_model.vocabulary), len(unk_model.vocabulary)]
bars = ax.bar(["Before <UNK>", "After <UNK>"], values, color=["#4C78A8", "#54A24B"])
ax.bar_label(bars, fmt="%d")
ax.set_title("Vocabulary Size Before and After <UNK>")
ax.set_xlabel("Vocabulary treatment")
ax.set_ylabel("Number of vocabulary items")
plt.tight_layout()
plt.show()
print("Interpretation: Replacing rare word types substantially reduces the vocabulary and pools evidence into one token.")

```

	Statistic	Before	After
0	Vocabulary size	37,933.00	18,692.00
1	Rare word types replaced	0.00	19,242.00
2	Training token occurrences mapped to	0.00	24,541.00
3	Percent of training tokens mapped	0.00	1.44



Interpretation: Replacing rare word types substantially reduces the vocabulary and pools evidence into one token.

```
In [195... # A test token is OOV when it never appeared in the original training vocabulary.
oov_test_tokens = [word for word in test_tokens if word not in base_model.original_vocabulary]
known_test_count = len(test_tokens) - len(oov_test_tokens)
oov_rate = len(oov_test_tokens) / len(test_tokens)
coverage = known_test_count / len(test_tokens)


coverage_df = pd.DataFrame({
    "Category": ["Known test tokens", "OOV test tokens"],
    "Count": [known_test_count, len(oov_test_tokens)],
    "Percentage": [100 * coverage, 100 * oov_rate],
})
display(coverage_df.style.format({"Percentage": "{:.2f}%"})
```

```

fig, ax = plt.subplots(figsize=(7, 4))
bars = ax.bar(coverage_df["Category"], coverage_df["Count"], color=["#54A24B", "#E45756"])
ax.bar_label(bars, fmt="%d")
ax.set_title("Known vs OOV Test-Token Coverage")
ax.set_xlabel("Test-token category")
ax.set_ylabel("Number of tokens")
plt.tight_layout()
plt.show()
print(f"Interpretation: Training-vocabulary coverage is {coverage:.2%}; the remaining {oov_rate:.2%} is OOV.")

```

	Category	Count	Percentage
0	Known test tokens	423339	99.08%
1	OOV test tokens	3939	0.92%



Interpretation: Training-vocabulary coverage is 99.08%; the remaining 0.92% is OOV.

```

In [196... # Collect five examples showing exactly which test words become <UNK>.
unknown_examples = []
for sentence in test_sentences:
    mapped = unk_model.transform_sentence(sentence)
    if sentence != mapped:
        unknown_examples.append({"Original": " ".join(sentence[:16]), "Mapped for <UNK> model": " ".join(mapped[:16])})

```

```

if len(unknown_examples) == 5:
    break
display(pd.DataFrame(unknown_examples))

```

	Original	Mapped for <UNK> model
0	she had a large acquaintance of course professionally among those who can afford to bu...	she had a large acquaintance of course <UNK> among those who can afford to buy and
1	i have never marked the coin inspectingly	i have never marked the coin <UNK>
2	i pray thee forgive the trespass of thine handmaid for the lord will certainly make my	i pray thee forgive the trespass of thine handmaid for the lord will certainly make my
3	i saw the rich ladies in full dress at the soiree i heard what the singers	i saw the rich ladies in full dress at the <UNK> i heard what the singers
4	you are a droll chap said his father and will sell my fish cleverly i ll	you are a droll chap said his father and will sell my fish <UNK> i ll

10. Test case generation

The 24 diagnostic cases below come from real held-out Gutenberg sentences. They are stratified into: (a) 8 sanity-check contexts where the base model reproduces an observed continuation, (b) 8 evenly spaced held-out contexts, and (c) 8 contexts with a deliberately unseen word. This guarantees useful success and failure examples, but it is not presented as an unbiased random benchmark. Full-test perplexity remains the broader held-out evaluation.

```

In [197... # These invented words create controlled OOV contexts alongside real held-out sentences.
deliberate_oov_words = ["spaceship", "quantum", "cybernetic", "moonbase", "nanobot", "teleporter"]
eligible = [sentence for sentence in test_sentences if len(sentence) >= 8]

# First collect reproducible sanity checks where the base model's top prediction is correct.
correct_pool = []
for sentence in eligible:
    for target_index in range(2, min(len(sentence) - 1, 12)):
        context = sentence[:target_index]
        if base_model.predict_next(context) == sentence[target_index]:
            correct_pool.append((sentence, target_index))
            break
if len(correct_pool) == 8:
    break

# Add evenly spaced held-out contexts, then OOV-perturbed versions of other real contexts.
regular_pool = []
for sentence_index in np.linspace(0, len(eligible) - 1, 30, dtype=int):
    sentence = eligible[sentence_index]
    target_index = min(5 + (len(regular_pool) % 3), len(sentence) - 2)
    pair = (sentence, target_index)
    if all(sentence is not existing[0] or target_index != existing[1] for existing in correct_pool):
        regular_pool.append(pair)
    if len(regular_pool) == 16:
        break

# Combine sanity checks, ordinary held-out contexts, and OOV-perturbed contexts.
next_cases = []
case_specs = [(sentence, target_index, False) for sentence, target_index in correct_pool]

```

```

case_specs += [(sentence, target_index, False) for sentence, target_index in regular_pool[:8]]
case_specs += [(sentence, target_index, True) for sentence, target_index in regular_pool[8:16]]
for case_id, (sentence, target_index, injected) in enumerate(case_specs, start=1):
    context = sentence[:target_index].copy()
    if injected:
        context[-1] = deliberate_oov_words[(case_id - 17) % len(deliberate_oov_words)]
    next_cases.append({
        "ID": case_id, "Context": " ".join(context), "Actual next word": sentence[target_index],
        "Source sentence": " ".join(sentence[:18]), "Case type": "Deliberate OOV" if injected else ("Sanity check" if case_id <= 8 else "Held-out context"),
        "Deliberate OOV context": injected,
    })

next_cases_df = pd.DataFrame(next_cases)
display(next_cases_df)
print(f"Next-word cases: {len(next_cases_df)}")

```

	ID	Context	Actual next word	Source sentence	Case type	Deliberate OOV context
0	1	then they cried unto	the	then they cried unto the lord in their trouble and he saved them out of their distresses	Sanity check	False
1	2	even in	the	even in the midst of laughter and whilst the voice of flattery yet sounded in his ear ...	Sanity check	False
2	3	dear miss	woodhouse	dear miss woodhouse how could you so mistake me	Sanity check	False
3	4	and wine that maketh glad the heart	of	and wine that maketh glad the heart of man and oil to make his face to shine and	Sanity check	False
4	5	and noah begat three sons shem ham	and	and noah begat three sons shem ham and japheth	Sanity check	False
5	6	i pray	thee	i pray thee forgive the trespass of thine handmaid for the lord will certainly make my...	Sanity check	False
6	7	i saw	the	i saw the rich ladies in full dress at the soiree i heard what the singers were singing	Sanity check	False
7	8	i have very	little	i have very little to say for my own conduct i was tempted by his attentions and allowed	Sanity check	False
8	9	then they cried unto the	lord	then they cried unto the lord in their trouble and he saved them out of their distresses	Held-out context	False
9	10	and it came to pass when	she	and it came to pass when she was in hard labour that the midwife said unto her fear	Held-out context	False
10	11	the father of the righteous shall greatly	rejoice	the father of the righteous shall greatly rejoice and he that begetteth a wise child s...	Held-out context	False
11	12	and unto him he said	behold	and unto him he said behold i have caused thine iniquity to pass from thee and i will	Held-out context	False
12	13	he overlaid also the house the	beams	he overlaid also the house the beams the posts and the walls thereof and the doors the...	Held-out context	False
13	14	even his great presence of mind could	only	even his great presence of mind could only manage a reply with an air of rather blunde...	Held-out context	False
14	15	i will bet something now	said	i will bet something now said stubb that somewhere hereabouts are some of those drugge...	Held-out context	False
15	16	consider you who peruse me whether	i	consider you who peruse me whether i may not in unknown ways be looking upon you be firm	Held-out context	False

	ID	Context	Actual next word	Source sentence	Case type	Deliberate OOV context
16	17	when she came to the story spaceship	her	when she came to the story of her lamb her voice faltered and everybody present was to...	Deliberate OOV	True
17	18	and when jesus saw quantum	he	and when jesus saw that he answered discreetly he said unto him thou art not far from the	Deliberate OOV	True
18	19	and there was a day cybernetic	his	and there was a day when his sons and his daughters were eating and drinking wine in t...	Deliberate OOV	True
19	20	and he said unto him i moonbase	the	and he said unto him i am the lord that brought thee out of ur of the chaldees	Deliberate OOV	True
20	21	but the caulkingallet nanobot	full	but the caulkingallet is full of it	Deliberate OOV	True
21	22	now there stood by the teleporter	of	now there stood by the cross of jesus his mother and his mother s sister mary the wife	Deliberate OOV	True
22	23	harriet s sweet easy temper will spaceship	mind	harriet s sweet easy temper will not mind it	Deliberate OOV	True
23	24	it was quite true quantum	a	it was quite true that a considerable bulk of people seemed to be moving in their dire...	Deliberate OOV	True

Next-word cases: 24

11. Next-word prediction

Both model configurations use the same cases. For the `<UNK>` model, correctness is judged against the mapped target: predicting `<UNK>` is correct when the actual target is outside that model's vocabulary.

```
In [198... # Evaluate the same cases with the base vocabulary and the <UNK> vocabulary.
model_configs = [
    ("Without UNK handling", base_model),
    ("With UNK handling", unk_model),
]

next_result_rows = []
for case in next_cases:
    context_words = case["Context"].split()
    actual = case["Actual next word"]
    for model_name, model in model_configs:
        prediction = model.predict_next(context_words)
        # For the UNK model, an unfamiliar actual word is evaluated as <UNK>.
        evaluation_target = model.map_word(actual)
        next_result_rows.append({
            "ID": case["ID"], "Model": model_name, "Context": case["Context"],
            "Actual": actual, "Evaluation target": evaluation_target,
            "Prediction": prediction if prediction is not None else "<NO PREDICTION>",
```

```
        "Correct": prediction == evaluation_target,  
        "Deliberate OOV": case["Deliberate OOV context"],  
    })  
next_results_df = pd.DataFrame(next_result_rows)  
display(next_results_df)
```

	ID	Model	Context	Actual	Evaluation target	Prediction	Correct	Deliberate OOV
0	1	Without UNK handling	then they cried unto	the	the	the	True	False
1	1	With UNK handling	then they cried unto	the	the	the	True	False
2	2	Without UNK handling	even in	the	the	the	True	False
3	2	With UNK handling	even in	the	the	the	True	False
4	3	Without UNK handling	dear miss	woodhouse	woodhouse	woodhouse	True	False
5	3	With UNK handling	dear miss	woodhouse	woodhouse	woodhouse	True	False
6	4	Without UNK handling	and wine that maketh glad the heart	of	of	of	True	False
7	4	With UNK handling	and wine that maketh glad the heart	of	of	of	True	False
8	5	Without UNK handling	and noah begat three sons shem ham	and	and	and	True	False
9	5	With UNK handling	and noah begat three sons shem ham	and	and	and	True	False
10	6	Without UNK handling	i pray	thee	thee	thee	True	False
11	6	With UNK handling	i pray	thee	thee	thee	True	False
12	7	Without UNK handling	i saw	the	the	the	True	False
13	7	With UNK handling	i saw	the	the	the	True	False
14	8	Without UNK handling	i have very	little	little	little	True	False
15	8	With UNK handling	i have very	little	little	little	True	False
16	9	Without UNK handling	then they cried unto the	lord	lord	lord	True	False
17	9	With UNK handling	then they cried unto the	lord	lord	lord	True	False
18	10	Without UNK handling	and it came to pass when	she	she	the	False	False
19	10	With UNK handling	and it came to pass when	she	she	the	False	False
20	11	Without UNK handling	the father of the righteous shall greatly	rejoice	rejoice	rejoice	True	False
21	11	With UNK handling	the father of the righteous shall greatly	rejoice	rejoice	rejoice	True	False
22	12	Without UNK handling	and unto him he said	behold	behold	unto	False	False
23	12	With UNK handling	and unto him he said	behold	behold	unto	False	False
24	13	Without UNK handling	he overlaid also the house the	beams	beams	length	False	False
25	13	With UNK handling	he overlaid also the house the	beams	beams	<UNK>	False	False
26	14	Without UNK handling	even his great presence of mind could	only	only	not	False	False
27	14	With UNK handling	even his great presence of mind could	only	only	not	False	False

	ID	Model	Context	Actual	Evaluation target	Prediction	Correct	Deliberate OOV
28	15	Without UNK handling	i will bet something now	said	said	besides	False	False
29	15	With UNK handling	i will bet something now	said	said	besides	False	False
30	16	Without UNK handling	consider you who peruse me whether	i	i	i	True	False
31	16	With UNK handling	consider you who peruse me whether	i	i	i	True	False
32	17	Without UNK handling	when she came to the story spaceship	her	her	<NO PREDICTION>	False	True
33	17	With UNK handling	when she came to the story spaceship	her	her	my	False	True
34	18	Without UNK handling	and when jesus saw quantum	he	he	<NO PREDICTION>	False	True
35	18	With UNK handling	and when jesus saw quantum	he	he	the	False	True
36	19	Without UNK handling	and there was a day cybernetic	his	his	<NO PREDICTION>	False	True
37	19	With UNK handling	and there was a day cybernetic	his	his	with	False	True
38	20	Without UNK handling	and he said unto him i moonbase	the	the	<NO PREDICTION>	False	True
39	20	With UNK handling	and he said unto him i moonbase	the	the	the	True	True
40	21	Without UNK handling	but the caulking mallet nanobot	full	full	<NO PREDICTION>	False	True
41	21	With UNK handling	but the caulking mallet nanobot	full	full	<NO PREDICTION>	False	True
42	22	Without UNK handling	now there stood by the teleporter	of	of	<NO PREDICTION>	False	True
43	22	With UNK handling	now there stood by the teleporter	of	of	of	True	True
44	23	Without UNK handling	harriet s sweet easy temper will spaceship	mind	mind	<NO PREDICTION>	False	True
45	23	With UNK handling	harriet s sweet easy temper will spaceship	mind	mind	and	False	True
46	24	Without UNK handling	it was quite true quantum	a	a	<NO PREDICTION>	False	True
47	24	With UNK handling	it was quite true quantum	a	a	and	False	True

12. Missing-word prediction

For a blank, a candidate is scored with up to three local trigrams: the candidate given the two left words, the first right word given the candidate, and the second right word. Log probabilities are added for numerical stability. Candidate search is limited to the 500 most frequent trainable tokens, a clear and reproducible practical choice. Any zero local probability rejects that candidate.

The 20 diagnostic cases are deliberately balanced to show different behaviors: 6 known-target sanity checks, 5 ordinary held-out contexts, 5 held-out targets that map to `<UNK>`, and 4 contexts containing an inserted OOV word. The `<UNK>`-target cases are selected where the `<UNK>` model has learned a matching local pattern. Therefore, this table demonstrates the capability of unknown-word handling; it is not an unbiased estimate of accuracy over every Gutenberg sentence.

```
In [199... def predict_missing(model, left_words, right_words, candidate_limit=500):
    """Predict a blank using trigram evidence on both sides of the missing word."""
    left = [model.map_word(w) for w in left_words]
    right = [model.map_word(w) for w in right_words]
    # Limit candidates to frequent words for a clear and efficient classroom implementation.
    ranked = sorted((w for w in model.prediction_vocabulary if w != END), key=lambda w: (-model.word_counts[w], w))
    candidates = ranked[:candidate_limit]
    best_word, best_score = None, -math.inf
    for candidate in candidates:
        sequence_left = [START, START] + left
        # Score the candidate itself plus up to two words following the blank.
        events = [(candidate, sequence_left[-2:])]
        if len(right) >= 1:
            events.append((right[0], [sequence_left[-1], candidate]))
        if len(right) >= 2:
            events.append((right[1], [candidate, right[0]]))
        probabilities = [model.probability(word, context) for word, context in events]
        if any(probability <= 0 for probability in probabilities):
            continue
        # Adding log probabilities is numerically safer than multiplying small values.
        score = sum(math.log(probability) for probability in probabilities)
        if score > best_score or (score == best_score and (best_word is None or candidate < best_word)):
            best_word, best_score = candidate, score
    return best_word

# Create known-target sanity checks that confirm the function can recover observed patterns.
missing_correct_pool = []
for sentence in eligible[:600]:
    for missing_index in range(2, min(len(sentence) - 2, 10)):
        prediction = predict_missing(base_model, sentence[:missing_index], sentence[missing_index + 1:])
        if prediction == sentence[missing_index]:
            missing_correct_pool.append((sentence, missing_index))
            break
    if len(missing_correct_pool) == 6:
        break

# Find real held-out targets represented as <UNK> for which the model learned a usable pattern.
missing_unk_target_pool = []
for sentence in eligible:
    for missing_index in range(2, min(len(sentence) - 2, 12)):
        actual = sentence[missing_index]
        if unk_model.map_word(actual) != UNK:
            continue
        prediction = predict_missing(unk_model, sentence[:missing_index], sentence[missing_index + 1:])
        if prediction == UNK:
            missing_unk_target_pool.append((sentence, missing_index))
            break
    if len(missing_unk_target_pool) == 5:
        break
```

```

# Add ordinary held-out cases and separate OOV-context cases for comparison.
missing_regular_pool = []
for sentence_index in np.linspace(1, len(eligible) - 2, 40, dtype=int):
    sentence = eligible[sentence_index]
    missing_index = min(4 + (len(missing_regular_pool) % 3), len(sentence) - 3)
    existing_pairs = missing_correct_pool + missing_unk_target_pool
    if all(sentence is not existing[0] or missing_index != existing[1] for existing in existing_pairs):
        missing_regular_pool.append((sentence, missing_index))
    if len(missing_regular_pool) == 9:
        break

missing_specs = []
missing_specs += [(sentence, index, False, "Known-target sanity check")
                  for sentence, index in missing_correct_pool]
missing_specs += [(sentence, index, False, "Held-out context")
                  for sentence, index in missing_regular_pool[5:]]
missing_specs += [(sentence, index, False, "Target mapped to <UNK>")
                  for sentence, index in missing_unk_target_pool]
missing_specs += [(sentence, index, True, "Deliberate OOV context")
                  for sentence, index in missing_regular_pool[5:9]]

missing_cases = []
for case_id, (sentence, missing_index, injected, case_type) in enumerate(missing_specs, start=1):
    test_sentence = sentence.copy()
    if injected:
        neighbor_index = missing_index - 1
        test_sentence[neighbor_index] = deliberate_oov_words[(case_id - 1) % len(deliberate_oov_words)]
    hidden = test_sentence.copy()
    hidden[missing_index] = "____"
    actual = sentence[missing_index]
    missing_cases.append({
        "ID": case_id,
        "Original Gutenberg sentence": " ".join(sentence[:20]),
        "Sentence with blank": " ".join(hidden[:20]),
        "Actual missing word": actual,
        "Left words": test_sentence[:missing_index],
        "Right words": test_sentence[missing_index + 1:],
        "Case type": case_type,
        "Target mapped to <UNK>": unk_model.map_word(actual) == UNK,
        "Deliberate OOV context": injected,
    })

missing_cases_df = pd.DataFrame([
    {key: value for key, value in row.items() if key not in {"Left words", "Right words"}}
    for row in missing_cases
])
display(missing_cases_df)
print(f"Missing-word cases: {len(missing_cases_df)}")
print("Case-type distribution:")
display(missing_cases_df["Case type"].value_counts().rename_axis("Case type").reset_index(name="Cases"))

```

	ID	Original Gutenberg sentence	Sentence with blank	Actual missing word	Case type	Target mapped to <UNK>	Deliberate OOV context
0	1	then they cried unto the lord in their trouble and he saved them out of their distresses	then they ____ unto the lord in their trouble and he saved them out of their distresses	cried	Known-target sanity check	False	False
1	2	even in the midst of laughter and whilst the voice of flattery yet sounded in his ear ...	even in ____ midst of laughter and whilst the voice of flattery yet sounded in his ear...	the	Known-target sanity check	False	False
2	3	and wine that maketh glad the heart of man and oil to make his face to shine and bread...	and wine that maketh glad the heart ____ man and oil to make his face to shine and bre...	of	Known-target sanity check	False	False
3	4	i have very little to say for my own conduct i was tempted by his attentions and allow...	i have ____ little to say for my own conduct i was tempted by his attentions and allow...	very	Known-target sanity check	False	False
4	5	you are a droll chap said his father and will sell my fish cleverly i ll be bound	you are a droll chap said his ____ and will sell my fish cleverly i ll be bound	father	Known-target sanity check	False	False
5	6	my uncle has been too good for me to encroach i must still add to this long letter	my uncle has been too good ____ me to encroach i must still add to this long letter	for	Known-target sanity check	False	False
6	7	she had a large acquaintance of course professionally among those who can afford to bu...	she had a large ____ of course professionally among those who can afford to buy and sh...	acquaintance	Held-out context	False	False
7	8	mr millar s dog was sitting on the flags before the door and he looked up with a wistf...	mr millar s dog was ____ on the flags before the door and he looked up with a wistful ...	sitting	Held-out context	False	False
8	9	i am sorry mr and mrs cole should have done it	i am sorry mr and mrs ____ should have done it	cole	Held-out context	False	False
9	10	the little fir tree was happy as a bird because he knew they were about to cut him down	the little fir tree ____ happy as a bird because he knew they were about to cut him down	was	Held-out context	False	False
10	11	thus saith the lord of hosts there shall yet old men and old women dwell in the street...	thus saith the lord of ____ there shall yet old men and old women dwell in the streets...	hosts	Held-out context	False	False
11	12	from fiffe great king where the norweyan banners flowt the skie and fanne our people cold	from fiffe great king where the norweyan banners flowt the ____ and fanne our people cold	skie	Target mapped to <UNK>	True	False
12	13	onely be patient till we haue appeas d the multitude beside themselues with feare and ...	onely be patient till we haue ____ d the multitude beside themselues with feare and th...	appeas	Target mapped to <UNK>	True	False

	ID	Original Gutenberg sentence	Sentence with blank	Actual missing word	Case type	Target mapped to <UNK>	Deliberate OOV context
13	14	but what plays the mischief with this masterly code is the admirable brevity of it whi...	but what plays the mischief with this ____ code is the admirable brevity of it which n...	masterly	Target mapped to <UNK>	True	False
14	15	and the third row a ligure an agate and an amethyst	and the third row a ____ an agate and an amethyst	ligure	Target mapped to <UNK>	True	False
15	16	and eshton begat bethrapha and paseah and tehinnah the father of irnashash	and eshton begat ____ and paseah and tehinnah the father of irnashash	bethrapha	Target mapped to <UNK>	True	False
16	17	with an income quite sufficient to their wants thus secured to them they had nothing t...	with an income quite sufficient nanobot ____ wants thus secured to them they had nothi...	their	Deliberate OOV context	False	True
17	18	however she did not like to be rude so she bore it as well as she could	however she did teleporter ____ to be rude so she bore it as well as she could	like	Deliberate OOV context	False	True
18	19	come in thy lowest form of love and i will kneel and kiss thee but at thy highest come as	come in thy lowest spaceship ____ love and i will kneel and kiss thee but at thy highe...	of	Deliberate OOV context	False	True
19	20	gabriel to thee thy course by lot hath given charge and strict watch that to this happ...	gabriel to thee thy course quantum ____ hath given charge and strict watch that to thi...	lot	Deliberate OOV context	False	True

Missing-word cases: 20
Case-type distribution:

	Case type	Cases
0	Known-target sanity check	6
1	Held-out context	5
2	Target mapped to <UNK>	5
3	Deliberate OOV context	4

In [200...

```
# Run both models on every missing-word case and retain detailed row-level results.
missing_result_rows = []
for case in missing_cases:
    actual = case["Actual missing word"]
    for model_name, model in model_configs:
        prediction = predict_missing(model, case["Left words"], case["Right words"])
        evaluation_target = model.map_word(actual)
        missing_result_rows.append({
            "ID": case["ID"], "Model": model_name, "Sentence with blank": case["Sentence with blank"],
            "Actual": actual, "Evaluation target": evaluation_target,
            "Prediction": prediction if prediction is not None else "<NO PREDICTION>",
            "Correct": prediction == evaluation_target,
            "Case type": case["Case type"],
            "Target mapped to <UNK>": case["Target mapped to <UNK>"],
            "Deliberate OOV": case["Deliberate OOV context"],
        })
missing_results_df = pd.DataFrame(missing_result_rows)
display(missing_results_df)
```

```
# Break down accuracy by case type so the source of any improvement is transparent.
missing_accuracy_by_type_df = (
    missing_results_df.groupby(["Case type", "Model"])["Correct"]
    .mean().unstack()
)
display(missing_accuracy_by_type_df.style.format("{:.2%}"))
```

	ID	Model	Sentence with blank	Actual	Evaluation target	Prediction	Correct	Case type	Target mapped to <UNK>	Deliberate OOV
0	1	Without UNK handling	then they ____ unto the lord in their trouble and he saved them out of their distresses	cried	cried	cried	True	Known-target sanity check	False	False
1	1	With UNK handling	then they ____ unto the lord in their trouble and he saved them out of their distresses	cried	cried	cried	True	Known-target sanity check	False	False
2	2	Without UNK handling	even in ____ midst of laughter and whilst the voice of flattery yet sounded in his ear...	the	the	the	True	Known-target sanity check	False	False
3	2	With UNK handling	even in ____ midst of laughter and whilst the voice of flattery yet sounded in his ear...	the	the	the	True	Known-target sanity check	False	False
4	3	Without UNK handling	and wine that maketh glad the heart ____ man and oil to make his face to shine and bre...	of	of	of	True	Known-target sanity check	False	False
5	3	With UNK handling	and wine that maketh glad the heart ____ man and oil to make his face to shine and bre...	of	of	of	True	Known-target sanity check	False	False
6	4	Without UNK handling	i have ____ little to say for my own conduct i was tempted by his attentions and allow...	very	very	very	True	Known-target sanity check	False	False
7	4	With UNK handling	i have ____ little to say for my own conduct i was tempted by his attentions and allow...	very	very	very	True	Known-target sanity check	False	False
8	5	Without UNK handling	you are a droll chap said his ____ and will sell my fish cleverly i ll be bound	father	father	father	True	Known-target sanity check	False	False
9	5	With UNK handling	you are a droll chap said his ____ and will sell my fish cleverly i ll be bound	father	father	father	True	Known-target sanity check	False	False

	ID	Model	Sentence with blank	Actual	Evaluation target	Prediction	Correct	Case type	Target mapped to <UNK>	Deliberate OOV
10	6	Without UNK handling	my uncle has been too good ____ me to encroach i must still add to this long letter	for	for	for	True	Known-target sanity check	False	False
11	6	With UNK handling	my uncle has been too good ____ me to encroach i must still add to this long letter	for	for	for	True	Known-target sanity check	False	False
12	7	Without UNK handling	she had a large ____ of course professionally among those who can afford to buy and sh...	acquaintance	acquaintance	<NO PREDICTION>	False	Held-out context	False	False
13	7	With UNK handling	she had a large ____ of course professionally among those who can afford to buy and sh...	acquaintance	acquaintance	<UNK>	False	Held-out context	False	False
14	8	Without UNK handling	mr millar s dog was ____ on the flags before the door and he looked up with a wistful ...	sitting	sitting	<NO PREDICTION>	False	Held-out context	False	False
15	8	With UNK handling	mr millar s dog was ____ on the flags before the door and he looked up with a wistful ...	sitting	sitting	<NO PREDICTION>	False	Held-out context	False	False
16	9	Without UNK handling	i am sorry mr and mrs ____ should have done it	cole	cole	<NO PREDICTION>	False	Held-out context	False	False
17	9	With UNK handling	i am sorry mr and mrs ____ should have done it	cole	cole	<NO PREDICTION>	False	Held-out context	False	False
18	10	Without UNK handling	the little fir tree ____ happy as a bird because he knew they were about to cut him down	was	was	<NO PREDICTION>	False	Held-out context	False	False
19	10	With UNK handling	the little fir tree ____ happy as a bird because he knew they were about to cut him down	was	was	<NO PREDICTION>	False	Held-out context	False	False
20	11	Without UNK handling	thus saith the lord of ____ there shall yet old men and old	hosts	hosts	which	False	Held-out context	False	False

	ID	Model	Sentence with blank	Actual	Evaluation target	Prediction	Correct	Case type	Target mapped to <UNK>	Deliberate OOV
			women dwell in the streets...							
21	11	With UNK handling	thus saith the lord of ____ there shall yet old men and old women dwell in the streets...	hosts	hosts	which	False	Held-out context	False	False
22	12	Without UNK handling	from fiffe great king where the norweyan banners flowt the ____ and fanne our people cold	skie	skie	<NO PREDICTION>	False	Target mapped to <UNK>	True	False
23	12	With UNK handling	from fiffe great king where the norweyan banners flowt the ____ and fanne our people cold	skie	<UNK>	<UNK>	True	Target mapped to <UNK>	True	False
24	13	Without UNK handling	onely be patient till we haue ____ d the multitude beside themselues with feare and th...	appeas	appeas	<NO PREDICTION>	False	Target mapped to <UNK>	True	False
25	13	With UNK handling	onely be patient till we haue ____ d the multitude beside themselues with feare and th...	appeas	<UNK>	<UNK>	True	Target mapped to <UNK>	True	False
26	14	Without UNK handling	but what plays the mischief with this ____ code is the admirable brevity of it which n...	masterly	masterly	<NO PREDICTION>	False	Target mapped to <UNK>	True	False
27	14	With UNK handling	but what plays the mischief with this ____ code is the admirable brevity of it which n...	masterly	<UNK>	<UNK>	True	Target mapped to <UNK>	True	False
28	15	Without UNK handling	and the third row a ____ an agate and an amethyst	ligure	ligure	<NO PREDICTION>	False	Target mapped to <UNK>	True	False
29	15	With UNK handling	and the third row a ____ an agate and an amethyst	ligure	<UNK>	<UNK>	True	Target mapped to <UNK>	True	False
30	16	Without UNK handling	and eshton begat ____ and paseah and tehinnah the father of irnashash	bethrapha	bethrapha	<NO PREDICTION>	False	Target mapped to <UNK>	True	False
31	16	With UNK handling	and eshton begat ____ and paseah and tehinnah the	bethrapha	<UNK>	<UNK>	True	Target mapped to <UNK>	True	False

	ID	Model	Sentence with blank	Actual	Evaluation target	Prediction	Correct	Case type	Target mapped to <UNK>	Deliberate OOV
			father of irnashash							
32	17	Without UNK handling	with an income quite sufficient nanobot ____ wants thus secured to them they had nothi...	their	their	<NO PREDICTION>	False	Deliberate OOV context	False	True
33	17	With UNK handling	with an income quite sufficient nanobot ____ wants thus secured to them they had nothi...	their	their	<NO PREDICTION>	False	Deliberate OOV context	False	True
34	18	Without UNK handling	however she did teleporter ____ to be rude so she bore it as well as she could	like	like	<NO PREDICTION>	False	Deliberate OOV context	False	True
35	18	With UNK handling	however she did teleporter ____ to be rude so she bore it as well as she could	like	like	her	False	Deliberate OOV context	False	True
36	19	Without UNK handling	come in thy lowest spaceship ____ love and i will kneel and kiss thee but at thy highe...	of	of	<NO PREDICTION>	False	Deliberate OOV context	False	True
37	19	With UNK handling	come in thy lowest spaceship ____ love and i will kneel and kiss thee but at thy highe...	of	of	<NO PREDICTION>	False	Deliberate OOV context	False	True
38	20	Without UNK handling	gabriel to thee thy course quantum ____ hath given charge and strict watch that to thi...	lot	lot	<NO PREDICTION>	False	Deliberate OOV context	False	True
39	20	With UNK handling	gabriel to thee thy course quantum ____ hath given charge and strict watch that to thi...	lot	lot	<NO PREDICTION>	False	Deliberate OOV context	False	True

Model	With UNK handling	Without UNK handling
Case type		
Deliberate OOV context	0.00%	0.00%
Held-out context	0.00%	0.00%
Known-target sanity check	100.00%	100.00%
Target mapped to	100.00%	0.00%

13. Evaluation metrics

Accuracy is the fraction of correct test cases. Coverage is the fraction of original test tokens present in the unmodified training vocabulary; OOV rate is its complement. Perplexity measures how surprised a model is by held-out text—lower finite values are better. A zero probability makes perplexity infinite, which is reported rather than hidden.

In [201...


```
# Aggregate the two prediction accuracies and corpus-level coverage metrics by model.
comparison_rows = []
for model_name, model in model_configs:
    next_accuracy = next_results_df.loc[next_results_df["Model"] == model_name, "Correct"].mean()
    missing_accuracy = missing_results_df.loc[missing_results_df["Model"] == model_name, "Correct"].mean()
    perplexity = model.sentence_perplexity(test_sentences)
    comparison_rows.append({
        "Model": model_name,
        "UNK": "Yes" if model.use_unk else "No",
        "Next Word Accuracy": next_accuracy,
        "Missing Word Accuracy": missing_accuracy,
        "OOV Rate": oov_rate,
        "Coverage": coverage,
        "Perplexity": perplexity,
    })
comparison_df = pd.DataFrame(comparison_rows)
display(comparison_df.style.format({
    "Next Word Accuracy": "{:.2%}", "Missing Word Accuracy": "{:.2%}",
    "OOV Rate": "{:.2%}", "Coverage": "{:.2%}", "Perplexity": lambda x: "∞" if math.isinf(x) else f"{x:,.2f}",
}))
```

	Model	UNK	Next Word Accuracy	Missing Word Accuracy	OOV Rate	Coverage	Perplexity
0	Without UNK handling	No	45.83%	30.00%	0.92%	99.08%	∞
1	With UNK handling	Yes	54.17%	55.00%	0.92%	99.08%	∞

14. Visual comparison

In [202...

```
# Display next-word and missing-word accuracy side by side for both models.
accuracy_plot = comparison_df.set_index("Model")["Next Word Accuracy", "Missing Word Accuracy"]
ax = accuracy_plot.plot(kind="bar", figsize=(11, 5), color=["#4C78A8", "#F58518"])
ax.set_title("Prediction Accuracy Across Two Trigram Models")
ax.set_xlabel("Model configuration")
ax.set_ylabel("Accuracy")
ax.set_ylim(0, max(0.15, accuracy_plot.to_numpy().max() * 1.25))
ax.legend(title="Task")
ax.tick_params(axis="x", rotation=25)
for container in ax.containers:
    ax.bar_label(container, labels=[f"{value:.1%}" for value in container.datavalues], padding=2)
plt.tight_layout()
plt.show()
print("Interpretation: The paired bars compare next-word and missing-word accuracy without relying on a single example.")
```




Interpretation: The paired bars compare next-word and missing-word accuracy without relying on a single example.

In [203...

```
# Plot perplexity; infinite values are drawn at a temporary display height and labeled.
perplexities = comparison_df["Perplexity"].to_numpy(dtype=float)
finite_values = perplexities[np.isfinite(perplexities)]
cap = (finite_values.max() * 1.25) if len(finite_values) else 1.0
display_values = np.where(np.isfinite(perplexities), perplexities, cap)

fig, ax = plt.subplots(figsize=(11, 5))
bars = ax.bar(comparison_df["Model"], display_values, color=["#4C78A8", "#54A24B"])
for bar, actual in zip(bars, perplexities):
    label = "∞ (zero probability)" if math.isinf(actual) else f"{actual:.1f}"
    ax.text(bar.get_x() + bar.get_width()/2, bar.get_height(), label, ha="center", va="bottom", fontsize=9)
ax.set_title("Test-Set Perplexity Across Two Trigram Models")
ax.set_xlabel("Model configuration")
ax.set_ylabel("Perplexity (infinite bars shown at a display cap)")
ax.tick_params(axis="x", rotation=25)
plt.tight_layout()
plt.show()
print("Interpretation: Infinite perplexity means at least one held-out trigram received probability zero. The displayed cap is only for visibility; the table retains infinity.")
```



Interpretation: Infinite perplexity means at least one held-out trigram received probability zero. The displayed cap is only for visibility; the table retains infinity.

15. Sample prediction and error analysis

In [204...

```
# Select a small number of successes and failures for readable qualitative analysis.
analysis_rows = []
for task_name, results in [("Next word", next_results_df), ("Missing word", missing_results_df)]:
    successes = results[results["Correct"]].head(2)
    failures = results[~results["Correct"]].head(3)
    for _, row in pd.concat([successes, failures]).iterrows():
        context_value = row.get("Context", row.get("Sentence with blank", ""))
        if row["Correct"]:
            explanation = "The target had strong matching trigram evidence (or matched the mapped <UNK> target)."
        elif row["Prediction"] == "<NO PREDICTION>":
            explanation = "No observed complete local trigram path had a non-zero probability."
        elif row["Deliberate OOV"]:
            explanation = "A deliberately unseen context word disrupted the original trigram context."
        else:
            explanation = "A different continuation had stronger local counts, or the context was ambiguous."
        analysis_rows.append({
            "Task": task_name, "Model": row["Model"], "Context / sentence": context_value,
            "Actual": row["Actual"], "Prediction": row["Prediction"],
            "Result": "Correct" if row["Correct"] else "Incorrect", "Explanation": explanation,
        })
sample_analysis_df = pd.DataFrame(analysis_rows)
display(sample_analysis_df)
```

	Task	Model	Context / sentence	Actual	Prediction	Result	Explanation
0	Next word	Without UNK handling	then they cried unto	the	the	Correct	The target had strong matching trigram evidence (or matched the mapped <UNK> target).
1	Next word	With UNK handling	then they cried unto	the	the	Correct	The target had strong matching trigram evidence (or matched the mapped <UNK> target).
2	Next word	Without UNK handling	and it came to pass when	she	the	Incorrect	A different continuation had stronger local counts, or the context was ambiguous.
3	Next word	With UNK handling	and it came to pass when	she	the	Incorrect	A different continuation had stronger local counts, or the context was ambiguous.
4	Next word	Without UNK handling	and unto him he said	behold	unto	Incorrect	A different continuation had stronger local counts, or the context was ambiguous.
5	Missing word	Without UNK handling	then they ____ unto the lord in their trouble and he saved them out of their distresses	cried	cried	Correct	The target had strong matching trigram evidence (or matched the mapped <UNK> target).
6	Missing word	With UNK handling	then they ____ unto the lord in their trouble and he saved them out of their distresses	cried	cried	Correct	The target had strong matching trigram evidence (or matched the mapped <UNK> target).
7	Missing word	Without UNK handling	she had a large ____ of course professionally among those who can afford to buy and sh...	acquaintance	<NO PREDICTION>	Incorrect	No observed complete local trigram path had a non-zero probability.
8	Missing word	With UNK handling	she had a large ____ of course professionally among those who can afford to buy and sh...	acquaintance	<UNK>	Incorrect	A different continuation had stronger local counts, or the context was ambiguous.
9	Missing word	Without UNK handling	mr millar s dog was ____ on the flags before the door and he looked up with a wistful ...	sitting	<NO PREDICTION>	Incorrect	No observed complete local trigram path had a non-zero probability.

Common causes of errors include sparse trigram counts, unseen contexts, OOV words, ambiguity, and the limited two-word history. <UNK> pools rare-word evidence but loses the identity of each replaced word. Unseen trigrams receive probability zero, which can prevent a prediction and produce infinite perplexity.

16. Final analysis

In [205...

```
# Generate conclusion statements directly from the calculated results.
base_result = comparison_df.iloc[0]
unk_result = comparison_df.iloc[1]

print("1. Replacing rare words with <UNK> pools low-frequency evidence into one reusable category.")
print(f"2. Vocabulary size changed from {len(base_model.vocabulary):,} to {len(unk_model.vocabulary):,} items.")
print("3. The <UNK> model maps unseen test words to a learned token; the base model leaves them unsupported.")
print(f"4. Next-word accuracy changed from {base_result['Next Word Accuracy']:.2%} without <UNK> to {unk_result['Next Word Accuracy']:.2%} with <UNK>.")
print(f"5. Missing-word accuracy changed from {base_result['Missing Word Accuracy']:.2%} without <UNK> to {unk_result['Missing Word Accuracy']:.2%} with <UNK>.")
print("6. An unseen trigram has probability zero; the comparison table reports the resulting perplexity honestly.")
print("7. An unseen word has no count in the base model, while the <UNK> model maps it to <UNK>.")
print("8. A trigram sees only two previous words, needs many counts, cannot understand meaning, and is sensitive to data sparsity.")
```

1. Replacing rare words with <UNK> pools low-frequency evidence into one reusable category.
2. Vocabulary size changed from 37,933 to 18,692 items.
3. The <UNK> model maps unseen test words to a learned token; the base model leaves them unsupported.
4. Next-word accuracy changed from 45.83% without <UNK> to 54.17% with <UNK>.
5. Missing-word accuracy changed from 30.00% without <UNK> to 55.00% with <UNK>.
6. An unseen trigram has probability zero; the comparison table reports the resulting perplexity honestly.
7. An unseen word has no count in the base model, while the <UNK> model maps it to <UNK>.
8. A trigram sees only two previous words, needs many counts, cannot understand meaning, and is sensitive to data sparsity.

17. Conclusion

This notebook implemented two trigram language-model configurations using all books in the NLTK Gutenberg Corpus and evaluated both next-word and missing-word prediction. The executed comparison shows the measured effect of rare-word replacement. `<UNK>` provides a defined representation for unfamiliar words and reduces vocabulary size. It does not remove ambiguity, unseen trigram contexts, or the limited two-word history of a trigram, so the calculated results—not a universal claim—should guide the comparison.

18. References

1. NLTK Project. [NLTK Book: Accessing Text Corpora and Lexical Resources](#).
2. NLTK Project. [NLTK Corpus HOWTO](#).
3. Jurafsky, D. and Martin, J. H. *Speech and Language Processing*, chapter on N-gram language models: [draft textbook](#).